



PROJECT LIONS

90 Days to Agentic Engineering

From AI skepticism to productive,
controlled agentic workflows

Thomas Schöne

Lead AI Architect

AGNTCon + MCPCon Europe 2026

My Turning Point

MY IMPLEMENTATION

Several

hours

AGENT IMPLEMENTATION

< 30

minutes

The same ticket.

The agent had no knowledge of my solution.

Very beginning of agents. No MCP, skills or anything like that.



“

**I spent five minutes looking at these
Als and immediately realized they were
nonsense.**

A visitor at an energy industry conference

Five minutes became the wrong test.

Who Masters a Tool in Five Minutes?

Excel

Photoshop

SAP

Jira

Git

IDE

Docker

Kubernetes

AutoCAD

Blender

Professional tools have to be learned. AI is no different.

Firefox: The Bug That Waited 20 Years

20 YEARS

A latent XSLT use-after-free bug

Human expertise

Defined the target, threat model and verification.

Agentic pipeline

Searched, tested and filtered at scale.

Reproducible proof

Produced a testcase—not just a plausible claim.

Mozilla fixed it in 2026 as part of an AI-assisted hardening pipeline.

KEY INSIGHT

**With clear direction and verification,
they turn expert methods into
faster, scalable work.**

The Team Behind the Journey

4

PEOPLE

Product owner + developers

Almost no prior AI experience

Occasional experiments
with ChatGPT

Starting Point: Day 1

Experience on one side. Uncertainty across the team.

● “AI hallucinates.”

● “I could do it faster myself.”

● “How do I know the code is correct?”

● “What may the AI see or change?”

Goal for Day 90

Not this:

More tools
Maximum autonomy
Replacing people


Instead:

Repeatable workflows
Visible verification
Sensible boundaries
Measurable value

Autonomy is earned trust.

Journey

How do we turn one developer's breakthrough into a workflow the whole team can trust?



Scale trust before you scale autonomy.



Make Failure Cheap

1

Small tasks

2

Visible results

3

Human approval

Early wins must be easy to verify—and cheap to reverse.

How Do We Stay in Control?

01

How do I get AI to do what I actually want?

02

How do I reliably verify what it has done?

Direction and verification—not magic.

UserPref: Requirements and Context

A small backend that keeps frontend preferences consistent across devices.

TICKET

What to store
Selected language
Widget configuration

PROJECT WIKI

Project architecture
Context for implementation

With the relevant context, implementation became straightforward.

Trust Started Inside the IDE



Adoption accelerated when AI stopped being a separate destination.

The Prototype Was Too Convincing

What I meant
“A two-minute
prototype.”

What the customer saw
“The requirement is
already implemented.”

If it looks finished, people assume it is finished.

When an Agent Is Convincingly Wrong

Plausible is not the same as correct.

Agent

“Task completed.”

≠


Reality

Tests, diff and requirements decide.


Trust comes from making failures reliably visible.

Vibe Code to Learn. Engineer to Ship.

DISCOVERY



DELIVERY



The prototype answers what to build—not how production should be built.

We had a way to check the work.

What could we safely delegate?

With clear boundaries and human review.

A chatbot answers.

An agent works toward an outcome.

What Is an Agent?

A model that directs its own process and tool use to achieve a goal.



The loop matters more than a single answer.

The Question Behind the Objection



**“By the time I’ve explained it to the AI,
I could have done it myself.”**

WHAT IS ACTUALLY MISSING?

- 01** Does the agent have the context it needs?
- 02** Is the goal specific enough?
- 03** Have we defined what “done” means?

You Are Not Explaining It for the Agent

CONTEXT IS SHARED INFRASTRUCTURE

It already has to exist for everyone who joins the work.

Tickets

What needs to change.

Wikis

How the system works.

Decisions

Why it works that way.


If missing context can be misread, fix the documentation process.

The agent exposes the gap. It did not create it.



How Do We Repeat What Works?

If it worked twice, we stopped treating it as a personal trick.



Repeatability turns a demo into engineering practice.

What Is a Skill?

A reusable package that teaches an agent how to perform a specific task.

Instructions



Resources



Scripts

The agent loads the right skill when the task needs it.

Skills Need a Trust Boundary

RISKS

Hidden instructions

Unsafe scripts

Excessive permissions

Data leakage

Unmaintained code

APPROVAL PATH

Need

Propose

Review

Internal repo

Approved use

Only skills from the internal repository may be used.

What Is MCP?

MCP

Model Context Protocol

An open standard that lets AI applications use tools and data from other systems.


Examples: read a wiki page or update a Jira ticket.

How Can the Agent Use Our Systems?

EXAMPLE TASK

Implement the ticket and update the project wiki.


MCP connects AI applications to external systems.



The skill describes the procedure. MCP connects the systems.

Documentation Joined the Development Loop

BEFORE



WITH AN AGENT



The documentation update stays connected to the ticket.

My Agent Also Does Timesheets

ME

09:00–11:30 — Project One.

Logged 2h 30m in Jira.

AI

ME

11:30–14:00 — Project Two.

A USEFUL BONUS

**No forms.
No timers.
No Friday
afternoon
archaeology.**

How Does the Agent Find Internal Knowledge?

Retrieval-Augmented Generation gives a model relevant external knowledge at runtime.




Which project decisions matter for this ticket?

A Typical RAG Flow

PREPARE KNOWLEDGE



ANSWER



The model is not retrained. The prompt receives better evidence.

RAG IS DEAD.

**Or is naïve “retrieve five chunks
and hope” RAG dying?**

The need for internal context is growing—not shrinking.

We Introduced RAG - Not a RAG Platform

WE DID NOT BUILD A SEPARATE PLATFORM

Ingestion → Chunking → Embeddings → Dedicated vector index

WE INTRODUCED

Retrieval-augmented workflows

Flagship model + agent + governed wiki access

RAG is the pattern. We did not need to own every layer.

RAG Retrieves. MCP Connects. The Agent Acts.

RAG

What internal evidence do we need?

MCP

Where can we find it—and what can we do?

Agent

Do we have enough context to continue?

The ticket brings knowledge, tools and verification into one workflow.

Scale Trust, Not Just Usage

More users were not enough.
We needed shared guardrails.

Trusted skills

Approved,
versioned and
owned.

Controlled tools

Scoped access
through MCP.

Reviews

Diffs, tests and
human approval.

Ownership

Someone maintains
the system.

Give Autonomy a Risk Budget

ROLLBACK

More freedom

Low impact · Easy rollback

More approval

High impact · Hard rollback

IMPACT →

Stop, Rework, or Standardize

STOP

No reliable value.

REWORK

Valuable, but unsafe or inconsistent.

STANDARDIZE

Repeatable and verifiable.

Not every experiment deserves to become infrastructure.

Three Assumptions We Got Wrong

01 · A better prompt is enough

Reliability came from the surrounding system.

02 · A convincing prototype explains itself

It needs an explicit label and expectation setting.

03 · Repetitive code is safest to generate freely

Repetition benefits most from constraints and validation.

Measure Outcomes, Not AI Activity

Time
saved

Rework
avoided

Adoption

Workflow
completion

More prompts are not a business outcome.

What I Would Do Differently

2

Integrate into the IDE earlier.

Label prototypes more clearly.

Create the trusted-skill path sooner.

After five minutes, AI can look like nonsense.

The real question is whether we can learn to direct it—and verify what it does.

Good agents take work off our hands.

Good engineering teams remain in control.



THANK YOU FOR YOUR ATTENTION

Questions?

Thomas Schöne · Lead AI Architect

