

I Was the Bottleneck, Not the Agent

Vincent YSMAL · Senior Software Engineer, Datadog



■ the agent's done: code · test · PR ■ my done: read · deploy · click · feedback

agent-1 code test PR **done** me read dep

Agent-1 says done. I pick it up and start reading.

■ the agent's done: code · test · PR ■ my done: read · deploy · click · feedback




Agent-2 starts, overtakes me, and is done. I'm still on agent-1.

■ the agent's done: code · test · PR ■ my done: read · deploy · click · feedback



Agents 3 and 4 land while I'm on agent-2.

■ the agent's done: code · test · PR ■ my done: read · deploy · click · feedback



All eight are done. I've finished two reviews. The rest is queue.

Oh. I have created a problem.



Eight very fast coworkers.
One extremely slow QA department.

Me.



The honest math


80%

Parallelized across eight agents

20%



Still mine, eight times over

A reasonable “done”



Until you’re running eight of them.

And then, the part nobody wrote down



Agents didn't invent this

Eng Done.

QA Where can I test it?

Eng Well, it works on my machine.

PM Is that available for customers?

Every job I've had. Four definitions of done, none of them written down.

The agent wasn't failing its definition of done. I had failed to define the one I actually needed.



THE TASK

Add a feature flag to hide the export button.

What came back



Code



Tests



Pull request


✓ Code

✓ Tests


✓ Pull request

Done?

Flag on, then refresh



✓ Button hidden



✗ Button is back

✓ Tests: still green

Two definitions of done

THE AGENT'S DONE

Implementation

Tests

Pull request

DONE

MY ACTUALLY-DONE

Understand

Deploy

Exercise the workflow

Inspect the evidence

Decide

Move it left

Implementation

Tests

Deploy

Exercise

Inspect

Pull request

DONE

The work didn't disappear. It moved to the left of "done."

My review

→

The agent's entry requirement

What I ask for

WHAT I USED TO WRITE

Add a feature flag to hide the
export button.

WHAT I WRITE NOW

What I ask for

WHAT I USED TO WRITE

Add a feature flag to hide the export button.

WHAT I WRITE NOW

The button is hidden when the flag is on.

What I ask for

WHAT I USED TO WRITE

Add a feature flag to hide the export button.

WHAT I WRITE NOW

The button is hidden when the flag is on.

It stays hidden after a refresh.

What I ask for

WHAT I USED TO WRITE

Add a feature flag to hide the export button.

WHAT I WRITE NOW

The button is hidden when the flag is on.

It stays hidden after a refresh.

When the flag is off, nothing else changes.

“Done” is the claims you need proven,
not the plan you’d have followed.



What I'd tell a new joiner

WEEK ONE, FIRST PR

"Did you test it?"

"Here's how you deploy it."

"The logs live here."

WHAT I GAVE THE AGENT


A task.

No definition of done.

And the expectation that it already knew the rest.

Nobody expects a new joiner to be omniscient. I expected it of the agent.

The closed loop



↩ and back to Act

First, a deployment skill



One capability, invoked whenever a task needs live validation.


Is the right thing actually running?

- ✓ Source state identified
- ✓ Build completed
- ✓ Artifacts synchronized
- ✓ Service restarted
- ✓ Health check passed

A deployment command is an action.

A verified deployment is a fact.

Observe



The workflow a person would actually use, run by the agent, not by me.

Also test the case I didn't ask about

EXPECTED CASE

Flag on → the button is gone

CONTROL CASE

Flag off → the button is still there

Without the control case, deleting the button passes every test.

The claim fails

✗ It stays hidden after a refresh



Button reappears on reload

Same bug as before.

Found inside the loop, not in my browser.

Compare against independent evidence

Compare



```
flag.export_hidden read=false  
source=default (no persisted value)  
trace 4f2c · page_load
```

Logs and traces, not the agent's own account of what it wrote.

One signal is never enough

USER SURFACE

Live URL

Visible behavior

Screenshot

RUNNING SYSTEM

Logs

Traces

Database state

AGENT SESSION


Recorded session

Tool calls

Tool results

A tool call shows the model asked. A tool result shows it ran.

Fix, redeploy, rerun the same scenario



The same scenario that broke, not a new, easier one.

Two ways out of the loop



Proven

Here is the evidence.



Blocked

Here is exactly what is stopping me.

~~Quietly tested something easier~~

Prove the behavior, or tell me what's
stopping you. Those are the only two doors.



The PR as a diary

Refactored the flag provider to read from the shared config loader. Added a `useExportFlag` hook and wired it through the toolbar container. Updated the export button component to accept a visibility prop. Renamed `exportEnabled` to `exportVisible` for consistency with the other flags. Also touched the config loader to memoize reads, since the provider was re-reading on every render. Added unit tests for the hook and snapshot tests for the toolbar. Updated three fixtures that referenced the old prop name. Removed a dead branch in the legacy toolbar that no longer compiles under the new prop type. Left the mobile toolbar alone for now — it uses a separate component tree and did not need changes. Rebased on main twice to pick up the config loader changes from #4412.

Everything it did. Nothing about whether it works.

The PR as a decision packet

CLAIM

The export button is hidden when the flag is on, and stays hidden after a refresh.

EVIDENCE

- 🔗 Live link at commit a3f19c2
- 📄 Scoped logs: flag read on reload
- 👁 Recording of the workflow

RISK

Touches the shared config loader. Other flags read from it.

UNKNOWN

Couldn't verify on mobile. No device in the test environment.

What I answer instead

01 Do I believe the claim?


02 Does the evidence support it?

03 Do I accept the risk?

04 Do I understand what's unknown?

I made “done” more expensive
for the agent.

And cheaper for me.



But I still find myself babysitting.



Filed against my system, not the agent

BUG I checked the same thing by hand, again

Component	My system, not the agent
Steps to reproduce	Hand a task to an agent. Wait. Ask the same follow-up question I asked last time.
Expected	The environment already knows to check this.
Severity	Recurring, once per agent, per task

Improve the environment, not the conversation

Unclear expectation

→ **Instruction**

Repeated workflow

→ **Skill**

Fragile manual operation

→ **Tool**

Recurring failure

→ **Regression test**

The environment remembers, so I don't have to.

Where an instruction lives

📄 **AGENTS.md** read at the start of every session

Definition of done

- Deployed to an isolated environment, commit verified live
- Behavior exercised in a browser: expected case and control case
- Evidence attached: live link, scoped logs, recording
- Anything unproven listed explicitly

Skills: deploy · browser · evidence

I describe the behaviors to validate. The repository provides the operating instructions.

Prompting each step isn't autonomy

IN THE CONVERSATION

- > please deploy this
- > now check it's ready
- > now look at the logs
- > now run it again

IN THE ENVIRONMENT

Standing instructions

Verification skill

Deploy + health tool

Regression tests

The loop is the default. I don't prompt it.

A very expensive command line.

agent-1 code test PR done me read deploy click feedback

agent-2 code test PR done me read deploy click feedback

agent-3 code test PR done me read deploy click feedback

agent-4 code test PR done me waiting read depl


agent-5 code test PR done me

agent-6 code test PR done me

agent-7 code test PR done me

agent-8 code test PR done me

My role moved



I stopped finishing the agent's work, and
started deciding whether I believed it.



Can I turn this babysitting into tooling?



FEEDBACK

Scan to rate this talk

Goes to me and the event organizers.

sfeedback.com/Jv5gnu

I Was the Bottleneck, Not the Agent
Vincent YSMAL, Datadog



I Was the Bottleneck, Not the Agent