

Building a sovereign AI governance stack with open source.

Roman Swoszowski

VP Product Development · Grape Up

How many AI systems
are you running right now?

4

11

40+

4

11


40+

Who is calling which model, with what data?


Are the guardrails the same everywhere?

What does each new integration open up?

What does this cost, and who pays?




“Maybe we should tell them it’s AI?”




Inventory · Access · Evidence




Who is calling which model?




What do we actually have?




Who is allowed to use what? (and how do they get it)




What happened in that conversation? (and can you show it in a year)



Are the rules the same everywhere?





Gateway

LiteLLM · Kong AI Gateway · Envoy AI Gateway · Apache APISIX · Bifrost
· IBM ContextForge · ToolHive · Docker MCP Gateway · MCPJungle · MCPX

Registry

MLflow · OpenMetadata · official MCP Registry · Kubeflow Hub

Access

Keycloak · Zitadel · Ory · OPA · Cerbos · OpenBao · MCP EMA


Observability

Langfuse · Arize Phoenix · Comet Opik · OpenLLMetry · OpenLIT · Laminar · SigNoz


Guardrails

NeMo Guardrails · Guardrails AI · Presidio · Llama Guard · ShieldGemma · Granite Guardian

How to assemble?




The rate of change
is itself changing.




automate draft triage answer
forecast summarise reconcile route
review onboard translate schedule

govern

Control





1 gateway and registry first

2 registered $\neq$ controlled

3 design for sovereignty

aibostr.com