



AGNTCON + MCPCON EUROPE · AMSTERDAM · 18 SEPTEMBER 2026

FROM ADVISORY TO AUTONOMOUS

A staged model for agent adoption,
and what happens when you skip a stage

THE SETUP

A wholesaler. Sales order processing. A story from my time at Lleverage.

Orders arrive by email, in every format a customer can invent. PDFs, spreadsheets, photos of handwritten notes. A team of four reads them, checks them, and keys them into the ERP. All day.

Communicative customer. Access to the mailbox and the ERP in days, not weeks. Eval set of 100+ real orders. Passed.

Working agent in three weeks. Into production. Everyone happy.

WHAT HAPPENED NEXT

Three weeks to build.
Three months before anyone
trusted it.

The operators found the wrong orders. Everything got checked by hand after that.

WHY THE GAP EXISTS

01

Probabilistic software

It does what it thinks you meant. People feel that.

02

Years of intuition

They want to see it survive the edge cases first.

03

Fear of replacement

No upside for the person inside the loop.

04

The 100% illusion

“We don’t make mistakes.” One AI error proves the point.

THE UNKNOWN UNKNOWN

**The agent doesn't say "not sure".
It says "done" and proceeds.**

100%

on an eval set of 100+ real
orders

~200

orders a day in production. The
evals cannot cover every case
that walks in.

0

ways to know which ones fell
outside the evals without a
human reading the original

Checks after the fact catch some of it: totals, dates, known customers. They cannot check completeness. Seven line items extracted. Were there eight? Only the original knows, and only a human is reading it.

THE COST OF A MISTAKE SETS THE BAR

FINANCE

A wrong invoice: thousands of euros

Tolerance for error is effectively zero. Every edge case matters. Every field gets verified. Full autonomy needs a mountain of evidence.

WHOLESALE

A wrong delivery address: tens of euros

Bad, but recoverable. The threshold for acceptable risk is higher, so the stages can move faster.

`cost per error × error rate × volume = the evidence you need before
each step`

Ask in week one. The answer sets the pace of everything after it.

THE FOUR STAGES

STAGE 01

Advisory

The agent observes, analyses, suggests. Humans do all the work. A second opinion they can ignore.

Sales orders: the agent reads the email and proposes the order. The operator keys it in themselves.

Goal: exposure

STAGE 02

Semi-autonomous

The agent drafts. A human approves before anything happens.

Sales orders: the agent creates the order in draft. The operator reviews, edits, releases.

Where most projects spend longest

STAGE 03

Enhanced autonomy

The agent executes routine cases. Humans own exceptions and edge cases.

Sales orders: clean orders booked automatically. Anything unusual queued for a human.

The agent has earned the routine

STAGE 04

Full autonomy

The agent runs end to end. Humans monitor and intervene rarely.

Sales orders: every order booked. A dashboard watched. Rules owned by a named person.

Earned, not given

Every transition needs evidence and buy-in. **From the operators, not just the sponsor.**

IN THE MODEL'S TERMS

We shipped at stage three.
The operators had never seen stages one and two.

- 01 Skipped stages one, two and four.
- 02 Stepped back to stage two. Not to zero.
- 03 Closed the loop: released orders compared with proposals, continuously.

You get one chance, maybe two.

THE GATES

STAGE	THE AGENT MAY	THE HUMAN DOES	MOVE ON WHEN
01 Advisory	Read the mailbox. Propose the order. No writes.	Keys the order in. Compares it with the proposal.	Operators agree with the proposal on more than 90% of a 100-order batch. The feedback loop is live and answered daily.
02 Semi- autonomous	Create the order in draft.	Reviews, edits, releases.	Edit rate below 5% over four weeks. Zero leadership- visible failures. Eval set expanded with every production edge case found.
03 Enhanced	Book clean orders. Queue the rest.	Works the exception queue. Spot-checks the clean ones.	Exception rate stable. Queue cleared within SLA. Spot- check accuracy holds above the cost-of-error threshold. Escalation path exercised, not just documented.
04 Full	Book every order.	Watches the dashboard. Owns the rules.	Monitoring in place. A named operator owns changes to the agent’s rules. Drift caught by the system before an operator notices it.

Thresholds scale with your cost of error. Finance moves slower than wholesale, on purpose.

EVERY DEPLOYMENT SINCE

01

**Stages are the plan,
not the recovery**

02

**Operators in
from the first call**

03

**Reaction
to every action**

The build didn't get faster. **The road to trust got shorter.**

WHERE THE PROTOCOL FITS

Autonomy is a permission set. Not a model setting.

STAGE which tools are in the agent's set. Reads only, then a draft write, then a booked write.

GATE what sits on the write call. A human release, then deterministic checks with a human queue, then monitoring.

PROTOCOL MCP, APIs, or a mix. The integration layer gives the agent doors. The trust ladder decides which ones it may walk through, and when.

LIVE IS NOT THE END

The business keeps moving after the build.

The agent keeps acting on the version that existed the day it shipped.

Changes to what the agent **believes** go through the same gate as what it **does**.

proposed rule

evidence

diff

confirmation

Rules change. Context changes. **Someone has to own what the agent knows and decide what is still true.**



The trust gap isn't a bug.
It makes you earn the right to
remove the human from the loop.

Instead of assuming it.



**Three weeks to build.
Three months to trust.
Plan for both.**

The essays behind this talk: fdehub.org